

Deep Learning

Winter Holiday 2022

LinkedIn

1/1/23

- Open jupyter notebooks using command prompt

Open folder, conda activate deeplearning, jupyter notebooks

- Deep learning: neural networks w/ 3 or more layers
 - Imitates human decision making and information processing
 - Starts w/ random initialization and works toward right values via trial and error
- Linear Regression
 - Dependent & independent variables, slope, intercept
 - $Y = ax + b$
- Logistic Regression
 - binary model
 - Relation btw 2 or more variables
 - Output 0/1
 - $Y, x, \text{slope, intercept, } f - \text{activation function}$
 - $Y = f(ax + b)$
 - F used to convert the continuous variable coming out of $ax + b$ into a boolean value
- Perceptron: unit for learning in artificial neural network
 - Represents an algo for supervised learning for binary classification
 - Like the cell in human brain // a cell in the neural network

- Based on logistic regression
- Replace slope w/ weight, w and intercept w/ bias, b
- Weights and biases are the parameters for a neural network
- Artificial neural networks (ANN) : networks of perceptrons
 - Perceptrons called nodes
 - Nodes organised as layers
 - Each node w/ its own weight. Biases and f_s
 - Each node connected to all nodes in next layer (w/ exceptions)
 - Working of ANN:
 - Inputs (independent var) sent from input layer
 - Passed onto nodes in hidden layer
 - Each node computes output based on its weights, biases and f
 - Node output passed onto next layer
- Training an ANN
 - Model is represented by parameters and hyperparameters
 - training : determining optimal " " to max accuracy
 - Inputs, weight and biases may be n-array
 - Process:
 - Use training data and create network architecture with intuition
 - Start w/ random values for wights and biases
 - Compute error in output
 - Adjust weights and biases to reduce errors
 - Also fine tune hyperparameters by adjusting layers, nose counts and others
 - Until error is an acceptable value
- The input layer
 - Vectors: ordered list of values
 - Used as inputs
 - A tuple

- Usually defines as NUmPy array
 - Represents feature variables
- Sample: an instance of a real world example (data set made of features (c) and samples (r)) and features individual attributes of a sample
- Input preprocessing: features need to be converted to numeric representations

Input Type	Preprocessing Needed
Numeric	Centering and scaling
Categorical	Integer encoding, one-hot encoding
Text	TF-IDF, embeddings
Image	Pixels – RGB representation
Speech	Time series of numbers

- E.g raw data -> centered and scaled -> transposed (switch x and y in table) [optional]
- Hidden layers
 - Typically 2n nodes
 - Neural network's architecture is defined by no of layers and nodes
 - Fully connected
 - Each node "learns" smt abt the feature-target relationship and this knowledge is persisted in its weights and biases
 - Inc nodes and layers Inc accuracy (not always true)
 - Architecture decided via experimentation
- Weights and biases
 - Represent trainable parameters in ANN
 - At a layer level, weights and bias are handled as RAs

Layer	Inputs	Nodes	Weights	Biases
HL 1	3	4	12	4
HL 2	4	5	20	5
HL 3	5	3	15	3
Output	3	2	6	2
Total			53	14

-
- 3 inputs, 4 nodes, therefore each node has 3 weight values = 12 weights overall. One bias value per node, so 4 biases
- Weights and biases are maintained as a matrix as well as inputs & outputs
- Activation func
 - Determines if a node will propagate info onto the next layer
 - Therefore, filters noise and normalizes output
 - Converts output to non linear value bc of matrix multiplication
 - Critical in learning patterns in the model

Popular Activation Functions

Activation Function	Output
Sigmoid	0 to 1
Tanh	-1 to +1
Rectified Linear Unit (ReLU)	0 if $x < 0$; x otherwise
Softmax	Vector of probabilities, with sum=1

- Choice depends on problem and experimentation.
- Output layer
 - Activation func maybe diff than the hidden layer
 - E.g. Softmax activation used for classification problems
 - Output may need further post-processing to convert to business values
 - Output layer size depends on problem
 - 1 for binary classification
 - N for n-class classification

- 1 for regression problems

TRAINING A NEURAL NETWORK

- Input Preprocessing
 - Input data is split into:
 - Training set: used to fit/determine the parameters
 - Validation set: used for model selection/ fine tuning
 - Test set: used to measure the final model performance
 - 80:10:10 usual split of data
 - Need to select values for layers and nodes in layers, activation function and hyperparameters
 - Initial selection based on intuition
 - Adjusted based on results
 - Weights and bias parameters need to be initialized
 - 0 initialization (not recommended)
 - Random initialization: values from a std normal distribution (mean =0, SD =1)
- Forward Propagation
 - \hat{y} (y w/ cap) (y hat) : prediction; y: actual
 - Send each sample thru neural network and obtain value of \hat{y}
 - Repeat for all samples and collect a set of \hat{y}
 - Compare values of \hat{y} to y to obtain error rates
- Error in prediction
 - Loss and cost function
 - Loss: measures prediction error for a single sample
 - Cost: measures error across a set of samples

Popular Cost Functions

Cost Functions	Applications
Mean Square Error (MSE)	Regression
Root Mean Square Error (RMSE)	Regression
Binary Cross Entropy	Binary classification
Categorical Cross Entropy	Multi-class classification

- - Measuring accuracy:
 - Send a set of samples through ANN and predict outcome
 - Estimate prediction error btw predicted outcome and expected outcome using a cost function
 - Use back propagation to adjust weights based on the error value
- Back Propagation
 - Opp. of forward propogation
 - Start from output layer
 - Compute delta value based on error found
 - Apply delta to adjust weights and biases in the layer
 - Derive new error value
 - Back propagate new error to previous layer and repeat
- Gradient descent
 - Process of repeating forward and backward propagations in order to reduce error and move closer to the desired model
 - Repeat the learning process
 - Forward propagation
 - Estimate error
 - Backward propagate
 - Adjust weights and biases
- Batches and Epoch help control the number of passes during the learning process

- Batches: set of samples sent through ANN in a single pass
 - batch size is a hyperparameter that defines the number of samples to work through before updating the internal model parameters.
 - Training data set can be divided into one or more batches
 - Training data is sent to the ANN one batch at a time
 - Cost estimates and parameters updates one batch at a time
 - 2 types:
 - Batch gradient descent: batch size = training set size
 - Mini-batch gradient descent: batch size < training set size
- Epoch: no. of times the entire training set is sent through the ANN
 - Epoch has one or more batches
 - Training process completes when all epoch is complete
 - Epoch sizes can be higher for better accuracy
- Training set size = 1000, batch size = 128, epoch = 50
 - Batches per epoch = $\text{ceil}(1000/128) = 8$ // ceiling value
 - The last batch will have fewer samples than 128
 - Total iterations (passes) through ANN = $8 \times 50 = 400$
 - Weights and biases updated 400 times
- Validation and testing
 - Validation: after each epoch and corresponding parameter updates, model can be used to predict for the validation data set
 - accuracy/loss can be measured and investigated
 - Evaluation: data set used to evaluate
- ANN model
 - Parameters: weights and biases
 - Hyperparameters: no. of layers, nodes in each layer, f , cost func, batch size, epoch

DEEP LEARNING EXAMPLE 1 - THE IRIS CLASSIFICATION PROBLEM

```
#Use a Label encoder to convert String to numeric values  
#for the target variable
```

```
from sklearn import preprocessing  
label_encoder = preprocessing.LabelEncoder()  
iris_data['Species'] = label_encoder.fit_transform(  
    iris_data['Species'])
```

-